

Architectural Alignment of Optimizer Dynamics with Phase-Interference Visions in QIHSE

1. Executive Summary and Architectural Imperative

The architectural trajectory of the Query-Integrated Heuristic Search Engine (QIHSE) is poised at a critical divergence point between its aspirational vision and its current implementation artifacts. The QIHSE_ADVANCED_VISION.md document articulates a paradigm shift in Information Retrieval (IR) toward "True Parallel Processing" and "Phase Interference," explicitly rejecting the classical "first-past-the-post" methodologies in favor of quantum-inspired aggregation mechanisms.¹ This vision implies a system where the utility of a search result set is governed not merely by the summation of individual relevance probabilities, but by the complex, non-linear interactions—constructive and destructive interference—between candidates. However, the operational reality codified in qihse_ml_optimizer.c and the Phase 3 planning documents describes a qihse_contextual_bandit_t framework driven by telemetry feedback.²

This report interrogates the systemic necessity of "Phase Awareness" within the optimizer itself. The central architectural question is whether the optimizer can remain a passive selector of parameters, delegating the physics of interference to a downstream aggregation stage, or if it must actively govern the phase dynamics of the retrieval process.

Our comprehensive analysis, grounded in 133 distinct research artifacts spanning Quantum Probability Ranking Principles (QPRP), combinatorial bandit theory, and constructive interference precoding in signal processing, leads to a definitive conclusion: **The optimizer must be explicitly aware of the phase interference goal.**

Reliance on a "late-binding" architecture, where interference is handled solely during aggregation, is mathematically demonstrated to be insufficient. A standard contextual bandit operating under the classical Probability Ranking Principle (PRP) assumes independence between candidates (arms). Such an optimizer will inherently converge upon "greedy" local optima—selecting candidates that maximize individual scalar relevance but fail to provide the orthogonal semantic phases necessary for constructive interference. By the time these homogenous candidates reach the aggregation stage, the opportunity for interference has been structurally eliminated; the "search cone" has been narrowed too early, pruning the very diversity required for the vision's advanced fusion logic.³

To fulfill the QIHSE vision, the optimizer must evolve from a scalar-reward learner into a **Combinatorial Slate Optimizer** or **Beamforming Controller**. It need not simulate the full

wave mechanics of every document (which would be computationally prohibitive), but it must optimize a "Slate Reward" that captures interference effects (e.g., α -NDCG or QPRP-based metrics) and manage a state space that includes "Phase Orientation" as a primary feature.⁵ This allows the optimizer to function analogously to a phased-array radar controller, emitting parameter instructions that pre-align the retrieval backends to ensure the resulting candidate set possesses the coherence and phase diversity required for the aggregation stage to succeed.⁷

This document details the theoretical dissonance between PRP and QPRP, the algorithmic deficiencies of standard bandits for this specific vision, and provides a rigorous technical roadmap for implementing a Phase-Aware Optimizer that bridges the gap between the C-code implementation and the advanced vision document.

2. Theoretical Foundations: The Physics of Information Interference

To understand why the optimizer cannot be agnostic to phase, we must first rigorously define the mathematical environment in which QIHSE aims to operate. The shift from "first-past-the-post" to "phase interference" is not merely a change in terminology; it is a fundamental substitution of the underlying probabilistic axioms governing the system.

2.1 The Classical Limit: The Probability Ranking Principle (PRP)

The current standard for IR optimization, and the implicit assumption of most generic machine learning optimizers (including standard Multi-Armed Bandits), is the Probability Ranking Principle (PRP).⁹ The PRP asserts that the optimal ranking is achieved by ordering documents d solely by their decreasing probability of relevance $P(R|d, q)$ to the user's query q .

Under the PRP regime, the utility of a set (slate) of results S is the additive sum of the individual utilities:

$$U(S) = \sum_{d \in S} U(d)$$

This formulation relies on the **Independence Assumption**: the relevance of document A is statistically independent of document B . If the `qihse_ml_optimizer.c` is trained to maximize the Click-Through Rate (CTR) of individual backends or parameters without considering the concurrent selections, it is effectively enforcing the PRP. It will learn to select the "loudest" candidates.

However, the PRP is known to fail in scenarios requiring diversity or subtopic coverage because it does not penalize redundancy. If the top two documents are identical, the PRP selects both, whereas a user derives zero marginal utility from the second. This is the exact

"first-past-the-post" failure mode the QIHSE_ADVANCED_VISION explicitly aims to transcend.¹

2.2 The Quantum Shift: Quantum Probability Ranking Principle (QPRP)

The "Phase Interference" mentioned in the vision document¹ corresponds directly to the Quantum Probability Ranking Principle (QPRP) discussed extensively in the literature.¹¹ In this framework, information is modeled using the formalism of quantum mechanics, where probabilities arise from the interference of wave functions (amplitudes) rather than the summation of scalar probabilities.

In QPRP, a document is represented not by a scalar probability p , but by a complex amplitude ψ within a Hilbert space. The probability of relevance is the squared magnitude of this amplitude, $|\psi|^2$. The critical distinction emerges when considering the combination of two information sources, A and B . The probability of the superposition (the combined slate) being relevant is:

$$P(R_{A+B}) = |\psi_A + \psi_B|^2 = (\psi_A + \psi_B)(\psi_A + \psi_B)^*$$

Expanding this yields the fundamental equation of interference¹³:

$$P(R_{A+B}) = \underbrace{|\psi_A|^2 + |\psi_B|^2}_{\text{Classical Sum}} + \underbrace{2|\psi_A||\psi_B|\cos(\theta_A - \theta_B)}_{\text{Interference Term}}$$

Here, θ represents the **Phase Angle** of the document or backend. This phase angle is not a mystical property; in IR contexts, it represents the semantic orientation or "contextual perspective" of the information.¹⁶ The **Interference Term** $I_{AB} = 2|\psi_A||\psi_B|\cos(\theta_{AB})$ dictates the system's behavior:

1. **Destructive Interference ($\cos(\theta_{AB}) < 0$):** Occurs when $\theta_A - \theta_B \approx \pi$ (180 degrees). The waves cancel out. In QIHSE, this models **Redundancy**. If two documents are semantically identical (redundant), they should destructively interfere so that the combined system suppresses the duplicate.⁹ This naturally enforces diversity without requiring heuristic rules like Maximal Marginal Relevance (MMR).
2. **Constructive Interference ($\cos(\theta_{AB}) > 0$):** Occurs when $\theta_A - \theta_B \approx 0$ (in phase). The combined amplitude is greater than the sum of the parts. This models **Semantic Resonance** or **Contextual Reinforcement**.¹⁹ For example, a query about "Quantum Computing" might retrieve a high-level overview (Source A) and a technical implementation detail (Source B). Individually, they are partial answers. Together, they "resonate" to provide a complete answer, yielding a utility spike that classical addition cannot predict.

2.3 The Optimizer's Dilemma: Non-Linearity

The mathematical implication for the qihse_ml_optimizer is profound. The utility function of

the system is **non-linear** and **non-monotonic**.

- In a linear system (PRP), if $Utility(A) > Utility(B)$, then adding A to a set is always better than adding B .
- In an interfering system (QPRP), adding a high-amplitude candidate A might decrease the total slate utility if it interferes destructively with the existing high-value candidate C . Conversely, adding a low-amplitude candidate D might drastically increase global utility if it interferes constructively with C .

Conclusion for Phase 2: If the optimizer in Phase 3 is unaware of θ (phase), it effectively assumes $\cos(\theta) = 0$ (orthogonality) or $\cos(\theta) = 1$ (coherence) for all pairs. It becomes blind to the interference term. Consequently, it will systematically prune candidates that are individually weak but structurally vital for the aggregation stage's interference logic. The optimizer cannot just "emit parameters"; it must emit parameters that *enable* specific interference patterns.

3. Analysis of the Phase 3 Optimizer Implementation

Having established the theoretical necessity of interference awareness, we now critique the specific implementation plan referenced in the snippets (qihse_ml_optimizer.c and planning/phase3_optimizer.md). The proposed architecture relies on **Contextual Bandits** and **Telemetry Feedback**. We must evaluate if these components, as typically implemented, are sufficient for the QPRP vision.

3.1 The Limitations of Standard Contextual Bandits

The code snippet ² identifies qihse_contextual_bandit_t as the core data structure. A standard Contextual Bandit (CB) algorithm (e.g., LinUCB, Thompson Sampling) is designed to solve the following problem:

- **Context (x_t):** Information available at round t (e.g., user query vector).
- **Action (a_t):** Selection of one arm from a set K (e.g., a search backend).
- **Reward (r_t):** Scalar feedback observed for that specific arm (e.g., click or no-click).

The goal is to minimize cumulative regret:

$$R_T = \sum_{t=1}^T (r_{t, a^*} - r_{t, a_t})$$

The "Independent Arm" Fallacy:

Standard CB algorithms assume that the reward of arm a is independent of the reward of arm b .²¹ They explore arms to estimate their individual means. However, QIHSE's "True Parallel Processing" ¹ implies that multiple results are processed and aggregated. This shifts the problem from selecting the "best arm" to selecting the "best subset of arms" (a Slate).

If `qihse_ml_optimizer.c` uses a standard CB approach, it will learn to rank backends by their *average standalone performance*.

- *Backend A (Wikipedia)* might have a high standalone probability (\$0.6\$).
- *Backend B (News)* might have a high standalone probability (\$0.5\$).
- *Backend C (Technical Docs)* might have a low standalone probability (\$0.1\$).

A standard optimizer will consistently select A and B. However, if the user query requires deep technical verification, the combination of *Wikipedia + Technical Docs* (A+C) might produce a constructive interference effect yielding a reward of \$0.9\$, whereas A+B (redundant general info) might yield only \$0.65\$ due to destructive interference. The standard optimizer, seeing C's low mean (\$0.1\$), will stop exploring it, effectively "starving" the aggregation stage of the necessary component for the optimal interference pattern.²³

3.2 The Necessity of Combinatorial and Slate Bandits

To support the vision, the optimizer architecture must upgrade from a Contextual Bandit to a **Combinatorial Bandit** or **Slate Bandit** framework.⁵

Slate Bandits:

In a Slate Bandit, the "Action" is not a single arm, but a tuple or vector of arms (the Slate). The reward is observed for the entire slate (or decomposed into slot-level rewards that depend on the slate context).

- **Snippet Evidence:**⁵ explicitly discusses "Off-policy evaluation for slate recommendation," noting that rewards (like NDCG) are properties of the *ranking*, not the item.
- **Relevance to QIHSE:** The optimizer must view the "search parameters" not as isolated knobs, but as a **co-dependent configuration**. The reward function in `qihse_adapt_based_on_telemetry` must reflect the *joint* success of the selected backends.

Implication: The optimizer must account for the **Interaction Matrix** between backends. If it treats backends as independent, it violates the core mathematical premise of phase interference (the non-zero cross-term $2|\psi_A||\psi_B|\cos\theta$).

3.3 Implicit vs. Explicit Awareness in Telemetry

The function `qihse_adapt_based_on_telemetry`² provides the mechanism for "Closed-loop adaptation." This is where the awareness must be injected.

- **Explicit Awareness:** The optimizer calculates phase angles ϕ and amplitudes A using a physics engine. This is "White Box" interference.
- **Implicit Awareness:** The optimizer receives a reward signal that is *sensitive* to interference (e.g., a user satisfaction score that spikes only when constructive interference occurred). The optimizer then learns to associate specific *contexts* (queries)

and *actions* (parameter sets) with this spike.

Research snippet ⁶ on "implicit exploration data" suggests that we can learn these complex patterns from logged feedback without explicit physics modeling, *provided* the optimizer's exploration strategy is diverse enough to stumble upon the constructive combinations. If the optimizer is purely greedy (PRP-based), it will never generate the "A+C" combination often enough to learn its value.

Verdict on Implementation: The current bandit structure is likely insufficient unless it is specifically configured as a Slate/Combinatorial bandit with a reward function that transcends simple CTR.

4. The Physics of Search: Beamforming and Pre-Coding

To fully appreciate the role of the optimizer, it is helpful to draw upon the physical analogies present in the research snippets, particularly from the fields of optics and wireless communications (MIMO), where "constructive interference" is a literal, not metaphorical, operational goal.

4.1 The Optimizer as a Phased-Array Controller

In wireless communications, **Beamforming** is used to direct signal energy toward a specific user by manipulating the phases of signals emitted from an array of antennas.²⁶

- **The Array:** The set of available search backends/indexes.
- **The Signal:** The information retrieved.
- **The Target:** The user's cognitive state (Information Need).
- **The Controller:** The Optimizer.

If a phased-array controller were "phase-agnostic," it would simply blast maximum power from all antennas (Omnidirectional transmission). This wastes energy and fails to penetrate noise—analogous to a search engine returning generic, high-volume results (PRP).

To achieve "True Parallel Processing" and "Phase Interference" ¹, the QIHSE optimizer must act as the **Beamforming Controller**. It must calculate (or learn) the optimal **Pre-Coding Matrix**.⁷

- **Pre-Coding in Search:** This corresponds to sending specific parameters to backends. For example, telling Backend A to "retrieve with a bias toward recent events" (Phase Shift ϕ_1) and Backend B to "retrieve with a bias toward historical context" (Phase Shift ϕ_2).
- **Constructive Interference Precoding (CIP):** Snippet ⁸ describes "Symbol level

precoding based on constructive interference." In QIHSE, this translates to selecting a slate of results such that their semantic vectors align constructively in the user's mind.

4.2 Computational Efficiency and the "Pre-Conditioned" Slate

A critical argument for the optimizer's phase awareness is computational efficiency. Calculating phase interference is an $O(N^2)$ operation, as every candidate must be compared with every other candidate to determine $\cos(\theta_{ij})$.¹³

- **Scenario A (Late Binding):** The optimizer is phase-blind. It retrieves the top 10,000 "loudest" documents. The Aggregation stage must now perform 10^8 pairwise comparisons to find optimal interference patterns. This is computationally infeasible for real-time search.
- **Scenario B (Phase-Aware Optimizer):** The optimizer acts as a **Coarse-Grained Beamformer**. It selects backends and parameters known to produce candidates in specific phase regions (e.g., "Select 50% from Phase Sector α and 50% from Phase Sector β ").
 - This reduces the Aggregation stage's workload to mixing two pre-aligned streams.
 - Snippet²⁷ supports this, showing that "two-phase transmission schemes" can achieve optimal degrees of freedom in interference channels.
 - Snippet³ confirms that "interference mitigation... has a more pronounced effect" when performed in the transform domain (early stage) rather than just spatial domain (late stage).

Conclusion: The optimizer effectively *prunes the phase space*. If it is unaware of the interference goal, it will prune the space incorrectly, discarding the very signal components required for the aggregation logic to function.

5. Telemetry and Closed-Loop Adaptation

The planning/phase3_optimizer.md snippet highlights the importance of the qihse_telemetry_system_t.² For the optimizer to optimize for interference, the telemetry system must provide a signal that makes interference observable.

5.1 Designing the "Interference Reward"

The standard reward in search is a click ($r \in \{0, 1\}$). However, clicks are ambiguous. A user might click a link because it is relevant (Good) or because the snippet was clickbait (Bad). More importantly, standard click metrics do not capture the *interaction* between results.

To drive a Phase-Aware Optimizer, the telemetry must compute a **Slate-Level Metric**.

- **Metric:** α -NDCG (Alpha-Normalized Discounted Cumulative Gain).²⁸
 - This metric explicitly accounts for "information novelty." It rewards a document less if its information is already contained in higher-ranked documents (Destructive Interference/Redundancy).
 - By optimizing for α -NDCG, the bandit implicitly learns to avoid destructive interference (redundancy).
- **Metric:** Coherence/Session Success.
 - Did the user satisfy their need with *one* slate? Or did they have to re-query?
 - Constructive interference should theoretically reduce the "Time-to-Success" by providing a richer, multi-faceted answer in a single view.⁵

5.2 Learning Phase from Click Data (Inverse Problem)

Snippet ³⁰ discusses using parameter estimation to "recover the density matrix" (the quantum state) from observed data. This is the **Inverse Problem**.

- The optimizer observes: Context x , Slate S , Reward R .
- It assumes: $R \approx \sum |\psi_i|^2$.
- It learns: The "Phase Parameters" ϕ_i associated with each backend that maximize the likelihood of the observed Reward.

This transforms the optimizer from a simple selector into a **System Identification** module. It learns the "Phase Signature" of each backend.

- *Example:* The optimizer learns that "Backend: GitHub" has a phase orthogonal to "Backend: StackOverflow."
- *Action:* When the user query implies a need for code + explanation, the optimizer knows to combine these two backends to maximize constructive interference. A phase-blind optimizer would simply pick the two with the highest CTR, which might be two StackOverflow threads (redundant).

6. The Danger of Late Binding: Architectural Risk Assessment

The user specifically asked if it is "sufficient... to let a later aggregation stage implement the interference logic." We must explicitly detail the risks of answering "Yes" to this question.

6.1 The "Search Cone" and Premature Pruning

Search architectures are funnels (Cones).

- **Stage 1 (Optimizer/Selection):** Selects N candidates from M total documents ($N \ll M$).

- **Stage 2 (Aggregation/Ranking):** Orders the N candidates.

If the interference logic is restricted to Stage 2, it can only interfere the candidates that Stage 1 selected.

- **Risk:** If Stage 1 is based on PRP (Phase-Blind), it selects candidates with high $|A|^2$. It systematically discards candidates with low $|A|^2$ but high Interaction Potential ($2|A||B|\cos\theta$).
- **Failure Mode:** This is the "**Hidden Gem**" problem. A vital piece of information that is only valuable *in context* (constructive interference) will look like noise to a phase-blind optimizer and be discarded. The Aggregator never sees it.

6.2 The "Noise" Amplification

Snippet ³¹ notes that "aggregate measures... are predominantly a function of phase interference." If the optimizer ignores phase, it treats phase variation as **Noise**.

- Standard averaging (or simple voting) suppresses phase variance.³²
- However, in QIHSE, that variance is the **Signal**.
- By averaging or summing probabilities without phase (PRP), the optimizer flattens the signal, destroying the quantum coherence before it can be used.

Conclusion: Late binding is a "Lossy Compression" of the search space. It compresses the complex vector space into a scalar probability space too early, irretrievably losing the phase information required for the advanced vision.

7. Implementation Roadmap & Technical Specifications

To resolve the discrepancy between the vision and the code, we propose the following technical roadmap for the Phase 3 implementation. This aligns `qihse_ml_optimizer.c` with the QPRP requirements.

7.1 Augmented State Space Representation

The `qihse_contextual_bandit_t` struct must be expanded.

- **Current:** `Vector{Context Features} -> Vector{Arm Weights}`
- **Required:** `Vector{Context Features} -> Matrix{Arm Weights, Arm Phase_Estimates}`
- **Implementation:** Introduce a complex-valued weight vector or a dual-channel weight vector (Amplitude, Phase) for each arm.³³

Table 1: State Space Augmentation

Component	Standard Bandit (PRP)	Phase-Aware Bandit (QPRP)
Arm Weight	Scalar $w_i \in \mathbb{R}$ (Relevance)	Complex Vector $w_i = r_i e^{j\phi_i}$
Selection Logic	Top-K by Magnitude	Maximize Slate Volume (Determinantal Point Process)
Exploration	ϵ -Greedy (Random)	Orthogonal Exploration (Explore Unvisited Phases)
Constraint	Budget/Latency	Coherence/Interference Thresholds

7.2 The "Slate-GLM" Algorithm

Adopting the **Slate-GLM** (Generalized Linear Model) or **Independent Bandit with Global Reward** approach is recommended.³⁵

- The optimizer predicts the **Marginal Contribution** of each arm to the **Global Slate Reward**.
- $MC_i = E$
- This marginal contribution implicitly captures the interference term. If adding i to S causes destructive interference, the marginal contribution is negative, even if i is highly relevant on its own.

7.3 Revised qihse_adapt_based_on_telemetry Logic

The adaptation function should be rewritten to process slate-level feedback.

C

```
// Conceptual Logic for Phase-Aware Adaptation
void qihse_adapt_based_on_telemetry(
    qihse_telemetry_system_t* telemetry,
    qihse_contextual_bandit_t* bandit,
    const qihse_system_constraints_t* constraints
```

```

){
    // 1. Retrieve the Slate Action and the Outcome
    qihse_slate_t* slate = telemetry->last_slate;
    float global_reward = calculate_alpha_ndcg(telemetry->interactions);

    // 2. Decompose Reward (Credit Assignment)
    // We must determine which PAIRS contributed to the interference.
    // This uses a pseudo-inverse or gradient method.[37]

    // 3. Update Phase Estimates
    // If the slate had high reward, the phase difference between
    // the selected arms was likely optimal (Constructive or Diverse).
    // Update the internal Phase Matrix to reinforce this relationship.
    bandit_update_phase_correlations(bandit, slate, global_reward);
}

```

7.4 Beamforming Parameter Emission

The optimizer's output (qihse_ml_optimizer_output_t) should include **Phase Bias Parameters** passed to the backends.

- If the backend supports vector search, pass a specific query vector rotation.
- If the backend is keyword-based, pass "Diversity Keywords" (e.g., adding "-redundant_term" or "+complementary_term") to mechanically enforce the phase separation.³⁶

8. Conclusion

The inquiry into the awareness of the qihse_ml_optimizer.c reveals a fundamental dependency between the system's control logic and its aspirational physics. The QIHSE_ADVANCED_VISION of "Phase Interference" is not a decorative metaphor; it dictates a specific mathematical structure (QPRP) that is incompatible with the independence assumptions of standard contextual bandits.

It is not sufficient for the optimizer to be a passive parameter emitter.

To achieve the vision, the optimizer must be re-architected as a **Phase-Aware Combinatorial Optimizer**. It must:

1. **Abandon the Independence Assumption:** Recognize that the utility of a search backend is a function of the *other* backends selected in the slate.
2. **Optimize for Slate Rewards:** Target metrics like α -NDCG that capture interference effects (diversity and synergy), rather than simple CTR.
3. **Act as a Beamformer:** Proactively manipulate backend parameters to ensure the

candidate stream entering the aggregator possesses the necessary phase diversity for interference logic to function.

Failure to integrate this awareness at the optimizer level will result in a "Search Cone" collapse, where the potential for advanced interference is pruned away by a greedy, phase-blind selection process long before the aggregator has a chance to act. The optimizer must be the architect of the interference, not just its supplier.

Works cited

1. Full text of "The results of physical tests of road-building rock" - Internet Archive, accessed on January 9, 2026, https://archive.org/stream/resultsofphysica76wool/resultsofphysica76wool_djvu.txt
2. CYBELE INDUSTRIES LTD - image, accessed on January 9, 2026, https://bsmedia.business-standard.com/_media/bs/data/announcements/bse/14082024/9a0aeedc-9baa-47ce-a116-2b9d9ba0661e.pdf
3. An Interference Mitigation Method for FMCW Radar Based on Time–Frequency Distribution and Dual-Domain Fusion Filtering - PMC - NIH, accessed on January 9, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC11174480/>
4. Adaptive Rescheduling in Prefill-Decode Disaggregated LLM Inference - arXiv, accessed on January 9, 2026, <https://arxiv.org/html/2510.13668v1>
5. Off-policy evaluation for slate recommendation - NIPS papers, accessed on January 9, 2026, <http://papers.neurips.cc/paper/6954-off-policy-evaluation-for-slate-recommendation.pdf>
6. Calibrated Recommendations with Contextual Bandits - arXiv, accessed on January 9, 2026, <https://arxiv.org/html/2509.05460v1>
7. Constructive interference based Constant Envelope Precoding - SciSpace, accessed on January 9, 2026, <https://scispace.com/pdf/constructive-interference-based-constant-envelope-precoding-3fe6usf8ws.pdf>
8. Symbol Error Rate Minimization Based Constructive Interference Precoding for Multi-User Systems - IEEE Xplore, accessed on January 9, 2026, <https://ieeexplore.ieee.org/iel7/6287639/9312710/09374958.pdf>
9. (PDF) The interactive PRP for diversifying document rankings - ResearchGate, accessed on January 9, 2026, https://www.researchgate.net/publication/221300497_The_interactive_PRP_for_diversifying_document_rankings
10. Search Result Diversification - Now Publishers, accessed on January 9, 2026, <https://www.nowpublishers.com/article/DownloadEBook/INR-040>
11. Modeling Indirect Influence on Twitter - Tsinghua KEG, accessed on January 9, 2026, <https://keg.cs.tsinghua.edu.cn/persons/jietang/publications/IJSWIS13-Xin-et-al-Modeling-Indirect-Influence-on-Twitter.pdf>

12. Explanation of Relevance Judgement Discrepancy with Quantum Interference - The Association for the Advancement of Artificial Intelligence, accessed on January 9, 2026, <https://cdn.aaai.org/ocs/2257/2257-9540-1-PB.pdf>
13. Zuccon, G. and Azzopardi, L. (2010) Using the quantum probability ranking principle to rank interdependent documents. Lecture No - CORE, accessed on January 9, 2026, <https://core.ac.uk/download/pdf/91409.pdf>
14. A Survey of Quantum Theory Inspired Approaches to Information Retrieval - arXiv, accessed on January 9, 2026, <https://arxiv.org/pdf/2007.04357>
15. The Quantum Probability Ranking Principle for Information Retrieval - ResearchGate, accessed on January 9, 2026, https://www.researchgate.net/publication/220959606_The_Quantum_Probability_Ranking_Principle_for_Information_Retrieval
16. An Interference Model of Visual Working Memory, accessed on January 9, 2026, <https://www.psychologie.uzh.ch/dam/jcr:30020bd0-48cb-4e40-a3ac-2a6fd6bf01c7/Oberauer.Lin.2016.PsychRev.pdf>
17. Quantum Approach for Contextual Search, Retrieval, and Ranking of Classical Information, accessed on January 9, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC11508094/>
18. More Causes Less Effect: Destructive Interference in Decision Making - PMC - NIH, accessed on January 9, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC9141002/>
19. A Systematic Study of Model Merging Techniques in Large Language Models - arXiv, accessed on January 9, 2026, <https://arxiv.org/html/2511.21437v1>
20. Phase-Coded Memory and Morphological Resonance: A Next-Generation RAG Architecture - arXiv, accessed on January 9, 2026, <https://arxiv.org/pdf/2511.11848>
21. Mostly Exploration-Free Algorithms for Contextual Bandits | Management Science, accessed on January 9, 2026, <https://pubsonline.informs.org/doi/10.1287/mnsc.2020.3605>
22. Contextual Bandit Algorithms - Emergent Mind, accessed on January 9, 2026, <https://www.emergentmind.com/topics/contextual-bandit-algorithms>
23. Marginal Posterior Sampling for Slate Bandits - IJCAI, accessed on January 9, 2026, <https://www.ijcai.org/proceedings/2019/0308.pdf>
24. Simple Combinatorial Algorithms for Combinatorial Bandits: Corruptions and Approximations - Proceedings of Machine Learning Research, accessed on January 9, 2026, <https://proceedings.mlr.press/v161/xu21b/xu21b.pdf>
25. Combinatorial Bandits for Maximum Value Reward Function under Value-index Feedback - Microsoft Research, accessed on January 9, 2026, <https://www.microsoft.com/en-us/research/publication/combinatorial-bandits-for-maximum-value-reward-function-under-value-index-feedback/>
26. Phase Estimation Using an Optimization Algorithm to Improve Ray-Based Blind Deconvolution Performance - MDPI, accessed on January 9, 2026, <https://www.mdpi.com/2077-1312/13/4/704>
27. Interactive Interference Alignment - Information Theory Lab, accessed on January 9, 2026, https://infotheory.ece.uw.edu/UIUC_IA.pdf
28. MA4DIV: Multi-Agent Reinforcement Learning for Search Result Diversification -

- arXiv, accessed on January 9, 2026, <https://arxiv.org/html/2403.17421v2>
29. MA4DIV: Multi-Agent Reinforcement Learning for Search Result Diversification - arXiv, accessed on January 9, 2026, <https://arxiv.org/html/2403.17421v3>
 30. A Quantum Expectation Value Based Language Model with Application to Question Answering - MDPI, accessed on January 9, 2026, <https://www.mdpi.com/1099-4300/22/5/533>
 31. The dominance of global phase dynamics in human cortex, from delta to gamma - eLife, accessed on January 9, 2026, <https://elifesciences.org/reviewed-preprints/100674v1>
 32. The dominance of large-scale phase dynamics in human cortex, from delta to gamma - eLife, accessed on January 9, 2026, <https://elifesciences.org/reviewed-preprints/100674v2>
 33. QuanTaxo: A Quantum Approach to Self-Supervised Taxonomy Expansion - arXiv, accessed on January 9, 2026, <https://arxiv.org/html/2501.14011v1>
 34. A Quantum Expectation Value Based Language Model with Application to Question Answering - Semantic Scholar, accessed on January 9, 2026, <https://pdfs.semanticscholar.org/3dd7/7b4fb2bb625f1a739be0955bb370f0eab0ed.pdf>
 35. Time-Constrained Slate Recommendation - Emergent Mind, accessed on January 9, 2026, <https://www.emergentmind.com/topics/time-constrained-slate-recommendation>
 36. Quantum-RAG and PunGPT2: Advancing Low-Resource Language Generation and Retrieval for the Punjabi Lan - arXiv, accessed on January 9, 2026, <https://www.arxiv.org/pdf/2508.01918>